

Raspberry Pi SDCard Burner for Building Cloud and Compute Clusters

Gregor von Laszewski, Anand Sriramulu, Sub Raizada,
Akshay Kowshik, Daivik Uggehalli Dayanand, Fugang Wang
Indiana University, Bloomington, IN 47401
laszewski@gmail.com

Abstract—To create a cluster from Raspberry Pi's one needs to either use a headless setup or burn a number of SDC cards. Typically after the burning of SDCards, we are faced with additional setup steps. However these steps can be simplified while assuring that the SDCard is modified after the burning with ssh enabled, a public key injected, a unique hostname assigned, and a network address specified. This command is naturally also important in case we need to re-burn a card in case a card would become bad or the OS on it for some reason corrupted. While attaching a multi-card USB writer it is possible to write multiple cards at the time one needs to switch cards into a single card writer.

I. INTRODUCTION

The Raspberry Pi provides to the community a cheap platform with the ability to expose Linux and other operating systems to the masses. Due to its cost point it is easy to buy a PI and experiment with it. As such this platform has been ideal to lower the entry barrier to advanced computing from the university level to highschool and middle school and even elementary school. However the PI has also been used by universities and even national labs. Due to its availability and is convenient accessibility it has become a staple of our educational pipeline.

Due to its price point the PI can also be used to build cheap clusters putting forward a hardware platform ideal for experimenting with issues such as networking and cluster management as educational tool. Many such efforts exist to use a PI as a cluster environment

TODO: add links here

There are a variety of administrative tasks that need to be conducted to allow such clusters to be put in place. This not only includes the creation of a physical space for the cluster, but also the ability to initialize such a cluster with software.

There are a number of different approaches to placing software for clusters on the Pi. In general we distinguish the following setups PXE boot and OS preloaded mode on the SD Cards.

II. BOOT CONFIGURATION METHODS

We review a number of configuration methods including setups known under the terms *headless*, *network booting*, and *booting from SDCards* and describe them in more detail next.

A. Headless

Under headless mode we understand that the Pi can be set up without the ability to have a monitor. The setup assumes that you have another computer from which you can log into the Raspberry. For this to work the PI and the computer must be on the same network. This could be a wired or wireless connection. The most common setup discussed in the community is the wireless setup. Here one needs to modify the `/etc/wpa_supplicant/wpa_supplicant.conf` file and enter the `ssid` and the `password` of the network (see Lst. 1). This can be accomplished by changing the file on the SDCard.

Listing 1 wpa_supplicant.conf

```
country=us
update_config=1
ctrl_interface=/var/run/wpa_supplicant

network={
    ssid="<Name of your WiFi>"
    psk="<Password for your WiFi>"
}
```

In addition we need to add the ability to remotely login with ssh. To do so we just have to place the file `ssh` in the boot folder.

More information about this method can be found at

- <https://www.raspberrypi.org/documentation/configuration/wireless/headless.md>

B. Network Booting or PXE Boot

The abbreviation PXE stands for Pre-boot eXecution Environment and is associated with the network booting

process of clients that boot from a server reachable via the network

In network booting we boot the PI from a server that contains all important OS information. One of the Pi's in the cluster (or other computer in the same network) needs to be set up as a DHCP/TFTP server.

This is typically a bit more complex, but can be scripted entirely so that you could convert a freshly installed PI into such a server.

More information about this is available at

- https://www.raspberrypi.org/documentation/hardware/raspberrypi/bootmodes/net_tutorial.md
- <https://www.raspberrypi.org/documentation/hardware/raspberrypi/bootmodes/net.md>

In addition to setting up the server, each Pi must be set up as a client. On each PI we need to install a number of tools and services that interact with the server. The boot process will take a significant amount of time and we need to be careful not to have too many Pi's in boot mode so the network is not overloaded.

C. Boot From a SDCard

The most common method to boot a PI is to have the operating system burned on the SDCard. In other methods we still have to have an SDCard or conduct modifications on the OS prior to us using an alternative method. Thus you will see that it is convenient to be able to write SDCards en-mass and make modifications on them to support the one or the other install method.

Hence if we had a convenient tool that allows us to not only burn the OS, but place important information on the SDCard before we boot, this will significantly reduce our setup time and also free the network from unnecessary traffic during the boot process of many PI's at once. This motivates our work and we deliver and describe our convenient command line tool to make such SDCards available through simple command interactions.

D. Other Solutions

Other solutions that we will expand upon in more detail in a future version of this paper include:

- PiServer: <https://www.raspberrypi.org/blog/piserver/>
- Boot from EEPROM (PI4): <https://www.raspberrypi.org/documentation/hardware/raspberrypi/booteprom.md>
- Boot from USB

III. INTEGRATE FROM HERE ON

To create a cluster from Raspberry Pi's one needs to either use a headless setup or burn a number of SDCards. Typically after the burning of SDCards, we are faced with additional setup steps. However these steps can be simplified while assuring that the SDCard is modified after the burning with ssh enabled, a public key injected, a unique hostname assigned, and a network address specified. This command is naturally also important in case we need to re-burn a card in case a card would become bad or the OS on it for some reason corrupted. While attaching a multi-card USB writer it is possible to write multiple cards at the time one needs to switch cards into a single card writer.

IV. OVERVIEW

cm-burn is a program to burn many SD cards for the preparation of building clusters with Raspberry Pi's. The program is developed in Python and is portable on Linux, Windows, and OSX. It allows users to create readily bootable SD cards that have the network configured, contain a public ssh key from your machine that you used to configure the cards. The unique feature is that you can burn multiple cards in a row. A sample command invocation looks like Lst. 2:

Listing 2 Command line invocation

```
cm-burn --name red[5-7] \
        --key ~/.ssh/id_rsa.pub \
        --ips 192.168.1.[5-7] \
        --image 2018-06-27-raspbian-stretch
```

This command creates 3 SD cards where the hostnames **red5**, **red6**, **red7** with the network addresses **192.168.1.5**, **192.168.1.6**, and **192.168.1.7**. The public key is added to the `authorized_keys` file of the `pi` user. The password login is automatically disabled and only the ssh key authentication is enabled.

V. PROCESS

The process of the burn is as follows.

1. start the program with the appropriate parameters the program will ask you to place an SD Card in the SD Card writer. Place it in
2. the specified image will be burned on the SD Card
3. next the SD Card will be mounted by the program and the appropriate modifications will be conducted.
4. after the modifications the SD Card will be unmounted
5. you will be asked to remove the card
6. if additional cards need to be burned, you will go to step 2.

In case a SD Card of a PI in the cluster goes bad, you can simply burn it again by providing the appropriate parameters, and just print the subset that are broken.

VI. SETTING UP A SINGLE LARGE CLUSTER WITH CM-BURN

cm-burn will setup a simple network on all cluster nodes configured. There are different models for networking configuration we could use. However we have decided for one that allows you to interface with your local Laptop to the cluster via Wifi. The setup is illustrated in Fig. 1

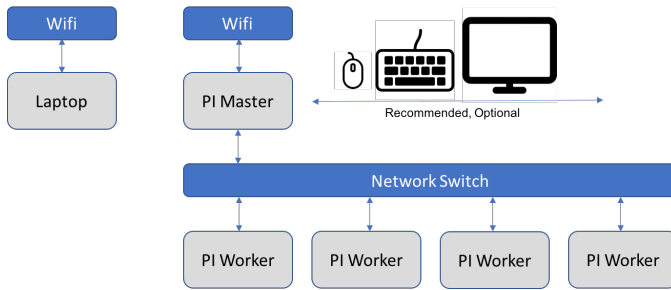


Figure 1: Network of a Raspberry Pi cluster

Figure: Networking

We assume that you have used **cm-burn** to create all SD cards for the Pi's. One of the Pi's is specially configured with the command

```
cm-burn --master red01
```

Manual page Lst. 3.

Listing 3 Creating a master

```
put the manual page here
```

The SD Card in the SD Card writer will be configured as a **master**. If the name does not match it will be configured as a worker. Only the **master** is connected with the Wifi network. All other nodes route the internet connection through the master node. As the **master** node is on the same Wifi network as the laptop you can login to the 'master' node and from there log into the workers. To simplify access you could even setup ssh tunneled connections from the Laptop via the master. But this is left up to you if you wish.

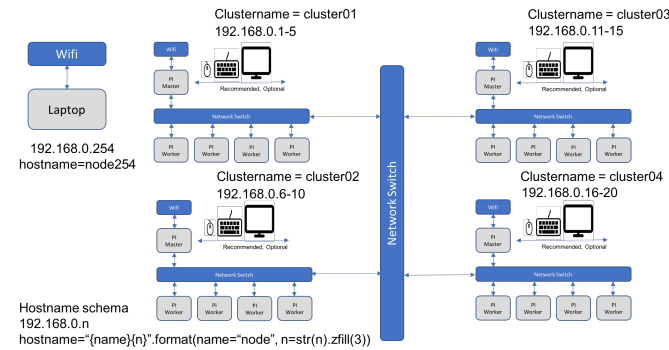
As a result you will be able to login on each of the machines and execute commands such as

```
sudo apt-get update
```

Certainly you can even have a much simpler setup by just attaching a keyboard, mouse and monitor/TV to your **master**. This will allow you to directly work on the master node, not needing any additional hardware.

A. Setting up a Cluster of Clusters with cm-burn

To integrate the clusters into a single network, we need a switch or combination of switches to which we connect the clusters. This is depicted in the Figure Cluster of Clusters



Each cluster is named cluster01-clusterNN. The hostnames are node followed by 3 zeros padded with the node number. There is a correlation between the cluster number and the node numbers in the following interval

a cluster has the nodes

$[(\text{clustername} - 1) * 5 + 1, (\text{clustername} - 1) * 5 + 5]$

For convenience we will be also enabling a cluster burn logic, that burns all images for a given cluster

```
cm-burn -workers=5 -name=cluster -nodes=nodes -id=3
```



B. Prerequisites

1) *Raspberry Pi*: We assume that you have set up a raspberry pi with the newest raspbian OS. We assume that you have changed the default password and can log into the pi.

We assume you have not done anything else to the OS.

The easiest way to duplicate the SD card is simply to clone it with the build in SD Card copier. This program can be found in the menu under Accessories.

Figure: SD Card Copier

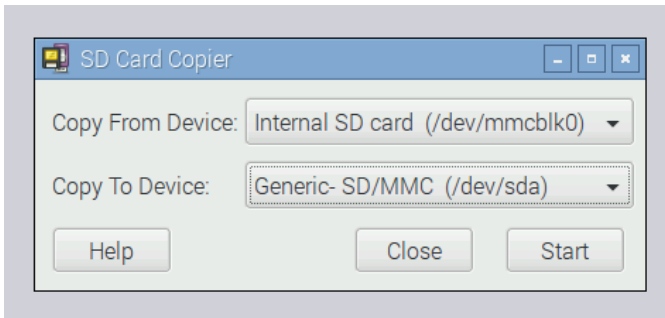


Figure 2: SD Card Copier

This program will copy the contents of the card plugged into the PI onto another one. The only thing you need is an USB SD Card writer. You can accept the defaults when the cards are plugged in which allow you to copy the Internal SD Card onto the other one. Just be careful that you do not overwrite your internal one. This feature can also be used to create backups of images that you have worked on and want to preserve.

Thus as you can see there is not much you need to do to prepare a PI to be used for burning the SD Card.

TODO: Python3

a) Card Burning from commandline:

- Insert card and find mmcblk0, e.g. no letter p in it for partition

```
sudo ls -ltr /dev/*
```

```
sudo dd bs=1M if=~/.cloudmesh/images/imagename.img of=mmcblk0 status=progress conv=fsync
```

2) OSX:

a) Card Burning:

On OSX a good program is to use etcher for burning the images on disk:

- <https://etcher.io/>

To access it from the commandline you can also use

- <https://etcher.io/cli/>

b) File System Management:

Unfortunately, the free versions of writing the ext file system are no longer supported on OSX. This means that as of writing of this document the best solution we found is to purchase and install extFS on the MacOS computer you use for burning the SD Cards. If you find an alternative, please let us know. (We tested ext4fuse, which unfortunately only supports read access, see Appendix)

To easily read and write ext file systems, please install extFS which can be downloaded from

- <https://www.paragon-software.com/home/extfs-mac/>

The purchase price of the software is \$39.95.

If you like to not spend any money we recommend that you conduct the burning on a raspberry pi.

TODO: PYTHON3 use pyenv

Tip: An alternative would be using virtualbox and using a virtual machine to avoid purchasing extFS.

C. Windows

a) Elevate permissions for Python.exe in Windows:

- Create a shortcut for python.exe
- Change the shortcut target into something like `C:\xxx\...\python.exe`
- Click “advance...” in the property panel of the shortcut, and click the option “run as administrator”

b) Executable needed to burn the image on SD Card::

Download CommandLineDiskImager from the following url

- <https://github.com/davidferguson/CommandLineDiskImager>

The above executable will be used by cm-burn script.

It's necessary to burn the raspbian image to the SD card with this executable manually or thru Etcher in order to continue with next step.

CommandLineDiskImager.exe

C:\Users\John\Downloads\raspbian.img G

c) File System Management:

Download the Open source ext3/4 file system driver for Windows installer from

- <http://www.ext2fsd.com/>
- Open Ext2fsd exe
- The burned image in the previous step in SD card will have 2 partition
- FAT32 partition will be assigned with the Drive letter - Boot Drive
- Assign Drive Letter for EXT4 (Right click on the EXT4, Assign letter. The drive letter will be used while running cm-burn) - Root Drive
- Setting Automount of this EXT4
- F3 or Tools->Ext2 Volume Management
- Check-> Automatically mount via Ext2Mgr
- The instructions above needed for the Ext2fsd to reserve the Drive Letters and any raspbian image burned to SD will be auto mounted to the specific reserved drive letters. These drive letters need to be specified while using cm-burn

D. Installation

1) *Install on your OS:* Once you have decided which Computer system (MacOS, Linux, or Windows) you like to use for using the cm-burn program you need to install it. The program is written in python3 which we assume you have installed and is your default python in your terminal.

To install cm-burn, please execute

```
git clone https://github.com/cloudmesh/cm-burn.git
cd cm-burn
pip install .
```

In future it will also be hosted on pypi and you will be able to install it with

```
pip install git+https://github.com/cloudmesh/cm-burn
```

To check if the program works please issue the command

```
cm-burn check install
```

It will check if you have installed all prerequisites and are able to run the command as on some OSes you must be in the sudo list to run it and access the SD card burner as well as mounting some file systems.

2) *Usage:*

a) *cmburn.yaml:*

```
cloudmesh:
  burn:
    image: None
```

b) *Manual page:*

1. git clone https://github.com/cloudmesh/cm-burn
2. cd cm-burn
3. python setup.py install
4. Copy the Rasperry PI images to be burned under
~/.cloudmesh/images

The manual page is as follows:

```
cm-burn -h
Cloudmesh Raspberry Pi Mass Image Burner.
```

Usage:

```
cm-burn create --group GROUP --names HOSTS --image IMAGE [--key=KEY] [--ips=IPS]
cm-burn gregor --group GROUP --names HOSTS --image IMAGE [--key=KEY] [--ips=IPS]
cm-burn ls
cm-burn rm IMAGE
cm-burn get [URL]
cm-burn update
cm-burn check install
cm-burn (-h | --help)
cm-burn --version
```

Options:

```
-h --help      Show this screen.
--version      Show version.
--key=KEY      the path of the public key [default: ~/.ssh/id_rsa.pub].
--ips=IPS      the ips in hostlist format
```

Location of the images to be stored for reuse:

```
~/.cloudmesh/images
~/.cloudmesh/inventory
```

Description:

```
cm-burn create [--image=IMAGE] [--group=GROUP] [--names=HOSTS]
               [--ips=IPS] [--key=PUBLICKEY] [--ssid=SSID]
               [--domain=DOMAIN]
               [--bootdrive=BOOTDRIVE] [--rootdrive=ROOTDRIVE]
               [-n --dry-run] [-i --interactive]

cm-burn ls [-ni]
cm-burn rm IMAGE [-ni]
cm-burn get [URL]
cm-burn update
cm-burn check install
cm-burn hostname [HOSTNAME] [-ni]
cm-burn ssh [PUBLICKEY] [-ni]
cm-burn ip IPADDRESS [--domain=DOMAIN] [-ni]
cm-burn wifi SSID [PASSWD] [-ni]
cm-burn info [-ni]
cm-burn image [--image=IMAGE] [--device=DEVICE] [-ni]

cm-burn (-h | --help)
cm-burn --version
```

Options:

```
-h --help      Show this screen.
-n --dry-run    Show output of commands but don't execute
-i --interactive Confirm each change before doing it
--version      Show version.
--key=KEY      the path of the public key [default: ~/.ssh/id_rsa.pub].
--ips=IPS      the IPs in hostlist format
--image=IMAGE  the image to be burned [default: 2018-07-16-raspbian-stretch.img]
```

Example:

```
cm-burn create --names red[000-010] ips --image raspbian.img
cmb-urn create --group g1 --names red[001-003] --key c:\ssh\id_rsa.pub
```

E. Appendix

1) *OSX ext4fuse:* Unfortunately ext4fuse only supports read access. To install it please use the following steps. However it will not allow you to use the cm-burn program. It may be useful for inspection of SD Cards

On OSX you will need brew and install osxfuse and ext4fuse

```
brew cask install osxfuse
brew install ext4fuse
```

To run it, your account must be in the sudoers list. Than you can do the following

```
mkdir linux
mkdir ~/.ssh
cp ../*.img 00.img
```

```
brew cask install osxfuse
brew install ext4fuse
hdiutil mount 00.img
```

This will return

```
/dev/disk3          FDisk_partition_scheme
/dev/disk3s1        Windows_FAT_32
/dev/disk3s2        Linux
```

We can now access the boot partition with

```
ls /Volumes/boot/
```

This partition is writable as it is not in ext format.

However to access the Linux partition in read only form we need to mount it with fuse

```
sudo mkdir /Volumes/Linux
sudo ext4fuse /dev/disk2s2 /Volumes/Linux -o allow_other
ext4fuse /dev/disk2s2 linux
less linux/etc/hosts
sudo umount /Volumes/Linux
```

2) *Activate SSH*: see method 3 in <https://www.raspberrypi.org/documentation/remote-access/ssh/>

Draft:

Set up ssh key on windows (use and document the ubuntu on windows thing)

you will have ~/.ssh/id_rsa.pub and ~/.ssh/id_rsa

copy the content of the file ~/.ssh/id_rsa.pub into ???/.ssh/authorized_keys ??? is the location of the admin user i think the username is pi

enable ssh on the other partition while creating the fike to activate ssh

3) *Hostname*: change /etc/hostname

4) *Activate Network*: see <https://www.raspberrypi.org/learning/networking-lessons/rpi-static-ip-address/>

5) *Change default password*: From the net (wrong method):

Mount the SD card, go into the file system, and edit /etc/passwd. Find the line starting with “pi” that begins like this:

```
pi:x:1000:1000...
```

Get rid of the x; leave the colons on either side. This will eliminate the need for a password.

You probably then want to create a new password by using the passwd command after you log in.

The right thing to do is to create a new hash and store it in place of x. not yet sure how that can be done a previous student from the class may have been aboe to do that Bertholt is firstname.

could this work? <https://unix.stackexchange.com/questions/81240/manually-generate-password-for-etc-shadow>

```
python3 -c "from getpass import getpass; from
crypt import *; p=getpass(); print('\n'+crypt(p,
METHOD_SHA512)) if p==getpass('Please repeat: ')
else print('\nFailed repeating.')"
```

F. Unmount Drives on Windows

RemoveDrive.exe needs to be downloaded to c:\Tools from the following path and to have the Administrator rights (Right Click on the exe -> Properties -> Compatibility Tab -> Run this program as an Administrator

- https://www.uwe-sieber.de/drivetools_e.html

See also

- <https://docs.microsoft.com/en-us/powershell/module/microsoft.powershell.management/remove-psdrive?view=powershell-6>

Gregor thinks that unmounting is much easier in an aelevated command prompt using

```
mountvol <Drive Letter>: /d
```

VII. LINKS

- https://github.com/cloudmesh-community/hid-sp18-419/blob/master/cluster/headless_setup.md
- <https://medium.com/@viveks3th/how-to-bootstrap-a-headless-raspberry-pi-with-a-mac-6eba3be20>
 - network setup is not good as it requires additional step, we want to preconfigure on sd card and plug in multiple pis at once not a single one.
- <https://github.com/cloudmesh/cloudmesh.pi/blob/dev/bin/cm-burn>
- http://www.microhowto.info/howto/mount_a_partition_located
- <http://www.janosgyerik.com/mounting-a-raspberry-pi-image-on-osx/>
- <https://github.com/Hitabis/pibakery>
- <http://osxdaily.com/2014/03/20/mount-ext-linux-file-system-mac/>
- <https://linuxconfig.org/how-to-mount-rasberry-pi-filesystem-image>
- <https://www.jeffgeerling.com/blogs/jeff-geerling/mounting-raspberry-pis-ext4-sd>
- <https://blog.hypriot.com/post/cloud-init-cloud-on-hypriot-x64/>
- <https://www.paragon-software.com/home/extfs-mac/>

VIII. OSX DURING BURNING

```
/dev/disk0 (internal):
```

```
#:          TYPE NAME
0:          GUID_partition_scheme
1:          EFI EFI
```

```
2:                Apple_APFS Container disk1                2.0 TB        disk0s2
```

```
/dev/disk1 (synthesized):
```

#:	TYPE NAME	SIZE	IDENTIFIER
0:	APFS Container Scheme -	+2.0 TB	disk1
	Physical Store disk0s2		
1:	APFS Volume Macintosh HD	811.4 GB	disk1s1
2:	APFS Volume Preboot	26.8 MB	disk1s2
3:	APFS Volume Recovery	519.0 MB	disk1s3
4:	APFS Volume VM	9.7 GB	disk1s4

```
/dev/disk2 (external, physical):
```

#:	TYPE NAME	SIZE	IDENTIFIER
0:	FDisk_partition_scheme	*31.9 GB	disk2

```
/dev/disk3 (external, physical):
```

#:	TYPE NAME	SIZE	IDENTIFIER
0:	FDisk_partition_scheme	*31.9 GB	disk3

Experiment DIY multiSDCard writer

We intend to experiment to build a multiSD card writer via USB. We will attempt to do this for OSX initially, therefore we like to order the following product

- USB Hub 3.0 Splitter, LYFNLOVE 7 Port USB Data

We will use multiple USB card readers (possibly just USB2 till we replacethem with USB3)

Than we will rewrite our program to attempt using the SDcard writers

Cloudmesh Raspberry Pi Mass Image Burner.

Usage:

```

cm-burn create [--image=IMAGE]
               [--group=GROUP]
               [--names=HOSTS]
               [--ips=IPS]
               [--key=PUBLICKEY]
               [--ssid=SSID]
               [--psk=PSK]
               [--bootdrive=BOOTDRIVE]
               [--rootdrive=ROOTDRIVE]
               [--domain=DOMAIN]
               [--bootdrive=BOOTDRIVE]
               [--rootdrive=ROOTDRIVE]
               [-n --dry-run]
               [-i --interactive]

cm-burn ls [-ni]
cm-burn rm IMAGE [-ni]
cm-burn get [URL]
cm-burn update
cm-burn check install
cm-burn hostname [HOSTNAME] [-ni]
cm-burn ssh [PUBLICKEY] [-ni]
cm-burn ip IPADDRESS [--domain=DOMAIN] [-ni]
cm-burn wifi SSID [PASSWD] [-ni]
cm-burn info [-ni]
cm-burn image [--image=IMAGE]
              [--device=DEVICE]
              [-ni]

cm-burn (-h | --help)
cm-burn --version

```

Options:

```

-h --help          Show this screen.
-n --dry-run       Show output of commands but don't execute them
-i --interactive    Confirm each change before doing it
--version          Show version.
--key=KEY          the path of the public key [default: ~/.ssh/id_rsa.pub].
--ips=IPS          the IPs in hostlist format
--image=IMAGE      the image [default: 2019-09-26-raspbian-buster.img]

```

Previously [default: 2018-06-27-raspbian-stretch.img]

Other images can be found at

<https://downloads.raspberrypi.org/raspbian/images/>

Files:

```

This is not fully thought through and needs to be documented
~/cloudmesh/images
~/cloudmesh/inventory
Location where the images will be stored for reuse

```

BUG:

```

bootdrive and rootdrive will be removed in a future release as they are self
discoverable

```

Description:

```

cm-burn
cm-burn create [--image=IMAGE]

```


REFERENCES

- [1] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [2] UCI Machine Learning Repository, “Spambase Data Set.” <https://archive.ics.uci.edu/ml/datasets/spambase>.
- [3] D. Graziotin, “How to write an ACM-styled conference paper using Markdown/Pandoc.” <https://ineed.coffee/4008/how-to-write-an-acm-styled-conference-paper-using-markdownpandoc>.
- [4] Misc., “Potter Ipsum.” <https://potteripsum.netlify.com>.